

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

In [4]:

```
mpg_df = pd.read_csv("C:\\Users\\pandey\\Downloads\\auto-mpg.csv")
```

In [5]:

```
mpg_df.head()
```

Out[5]: mpg cylinders displacement horsepower weight acceleration model year origin car name
0 18.0 8 307.0 130 3504 12.0 70 1 chevrolet chevelle malibu
1 15.0 8 350.0 165 3693 11.5 70 1 buick skylark
2 18.0 8 318.0 150 3436 11.0 70 1 plymouth satellite
3 16.0 8 304.0 150 3433 12.0 70 1 amc rebel sst
4 17.0 8 302.0 140 3449 10.5 70 1 ford torino

```
mpg_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   mpg              398 non-null    float64
1   cylinders        398 non-null    int64
2   displacement     398 non-null    float64
3   horsepower       398 non-null    object
4   weight           398 non-null    int64
5   acceleration     398 non-null    float64
6   model year      398 non-null    int64
7   origin           398 non-null    int64
8   car name        398 non-null    object
dtypes: float64(3), int64(4), object(2)
memory usage: 28.1+ KB
```

In [8]:

```
mpg_df.duplicated().sum()
```

Out[8]:

0

In [9]:

```
mpg_df.nunique()
```

Out[9]:

```
mpg          129
cylinders      5
displacement  82
horsepower    94
weight       351
acceleration  95
model year    13
origin        3
car name     305
dtype: int64
```

horsepower column have inconsistant data type.¶ In [10]:

```
mpg_df.horsepower.unique()
```

Out[10]:

```
array(['130', '165', '150', '140', '198', '220', '215', '225', '190',
      '170', '160', '95', '97', '85', '88', '46', '87', '90', '113',
      '200', '210', '193', '?', '100', '105', '175', '153', '180', '110',
      '72', '86', '70', '76', '65', '69', '60', '80', '54', '208', '155',
      '112', '92', '145', '137', '158', '167', '94', '107', '230', '49',
      '75', '91', '122', '67', '83', '78', '52', '61', '93', '148',
      '129', '96', '71', '98', '115', '53', '81', '79', '120', '152',
      '102', '108', '68', '58', '149', '89', '63', '48', '66', '139',
      '103', '125', '133', '138', '135', '142', '77', '62', '132', '84',
      '64', '74', '116', '82'], dtype=object)
```

In [11]:

```
(mpg_df['horsepower'] == '?').sum()
```

Out[11]:

6

In [12]:

```
mpg_df = mpg_df[mpg_df.horsepower != '?']
```

```
# confirming changes
(mpg_df.horsepower == '?').sum()
```

Out[12]:

0

Our dataset has a total of 398 records and 9 columns. We have no NaNs in our dataset nor duplicated rows. horsepower column have inconsistent data type that needs to be handled and casted to int. origin would need to be parsed and casted into a categorical datatype. No columns would need to be dropped.

In [13]:

```
mpg_df.horsepower = mpg_df.horsepower.astype(int)
```

```
# confirming changes
mpg_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 392 entries, 0 to 397
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   mpg             392 non-null    float64
 1   cylinders       392 non-null    int64
 2   displacement    392 non-null    float64
 3   horsepower      392 non-null    int32
 4   weight         392 non-null    int64
 5   acceleration    392 non-null    float64
 6   model year     392 non-null    int64
 7   origin          392 non-null    int64
 8   car name       392 non-null    object
dtypes: float64(3), int32(1), int64(4), object(1)
memory usage: 29.1+ KB
```

We'd map origin column according to the provided description (1 -> USA, 2 -> Europe, 3 -> Asia), and cast its datatype as category.

In [14]:

```
mpg_df.loc[mpg_df.origin == 1, 'origin'] = 'USA'
mpg_df.loc[mpg_df.origin == 2, 'origin'] = 'Europe'
mpg_df.loc[mpg_df.origin == 3, 'origin'] = 'Asia'

# casting origin column into category
```

```
mpg_df['origin'] = mpg_df['origin'].astype('category')

# validating changes
mpg_df['origin'].dtype

C:\Users\pandey\AppData\Local\Temp\ipykernel_10128\3030830440.py:1: FutureWarning
  mpg_df.loc[mpg_df.origin == 1, 'origin'] = 'USA'
```

Out[14]:

```
CategoricalDtype(categories=['Asia', 'Europe', 'USA'], ordered=False, categories_dtype=object)
```

CategoricalDtype(categories=['Asia', 'Europe', 'USA'], ordered=False) Now, our dataframe is tidy and clean, and we are ready to move into visualizing it to get some meaningful insights!

In [15]:

```
mpg_df.describe()
```

```
Out[15]: mpg cylinders displacement horsepower weight acceleration model year count 392.000000
392.000000 392.000000 392.000000 392.000000 392.000000 392.000000 mean 23.445918 5.471939
194.411990 104.469388 2977.584184 15.541327 75.979592 std 7.805007 1.705783 104.644004
38.491160 849.402560 2.758864 3.683737 min 9.000000 3.000000 68.000000 46.000000 1613.000000
8.000000 70.000000 25% 17.000000 4.000000 105.000000 75.000000 2225.250000 13.775000
73.000000 50% 22.750000 4.000000 151.000000 93.500000 2803.500000 15.500000 76.000000 75%
29.000000 8.000000 275.750000 126.000000 3614.750000 17.025000 79.000000 max 46.600000
8.000000 455.000000 230.000000 5140.000000 24.800000 82.000000 In [17]:
```

```
plt.figure(figsize=(10,5))
ax = sns.countplot(x = 'model year', data = mpg_df, color = '#4287f5')
ax.bar_label(ax.containers[0], label_type='edge')
plt.title("model year distribution", fontsize = 20)
plt.xlabel("model year", fontsize = 15)
plt.ylabel("cars count", fontsize = 15)
plt.show()
```

In [19]:

```
plt.figure(figsize=(10,5))
plt.title("model year against mpg", fontsize = 20)
plt.xlabel("model year", fontsize = 15)
plt.ylabel("mpg", fontsize = 15)
sns.lineplot(x = 'model year', y = 'mpg', data = mpg_df);
```

From the above visualizations, we we may totice the following:

Our dataset contains info about cars from 1970 to 1982.

Most of the cars are produced in 1973. As years pass after 1973, there has been a noticable increase in mpg.

In [20]:

```
plt.figure(figsize=(10,5))
ax = sns.histplot(x = 'mpg', data = mpg_df, color = '#4287f5')
ax.bar_label(ax.containers[0], label_type='edge')
plt.title("mpg count distribution", fontsize = 20)
plt.xlabel("mpg", fontsize = 15)
plt.ylabel("cars count", fontsize = 15)
plt.show()
```

In [21]:

```
plt.figure(figsize=(10,5))
ax = sns.countplot(x = 'cylinders', data = mpg_df, color = '#4287f5')
ax.bar_label(ax.containers[0], label_type='edge')
plt.title("cylinder distribution", fontsize = 20)
```

```
plt.xlabel("cylinder", fontsize = 15)
plt.ylabel("cars count", fontsize = 15)
plt.show()
```

In [22]:

```
plt.figure(figsize=(10,5))
sns.barplot(x = 'cylinders', y = 'mpg', data = mpg_df, color = '#4287f5')
plt.title("clynders against mpg", fontsize = 20)
plt.xlabel("clynders", fontsize = 15)
plt.ylabel("mpg", fontsize = 15)
plt.show()
```

In [23]:

```
plt.figure(figsize=(10,5))
sns.lineplot(x = 'cylinders', y = 'horsepower', data = mpg_df, color = '#4287f5')
plt.title("clynders against horsepower", fontsize = 20)
plt.xlabel("clynders", fontsize = 15)
plt.ylabel("horsepower", fontsize = 15)
plt.show()
```

In [24]:

```
plt.figure(figsize=(10,5))
sns.scatterplot(x = 'mpg', y = 'horsepower', data = mpg_df)
plt.title("clynders against mpg", fontsize = 20)
plt.xlabel("mpg", fontsize = 15)
plt.ylabel("cylenders", fontsize = 15)
plt.show()
```

From the above visuals, we can notice that:

As cylinders in the engine increases above 4, MPG decreases. As cylinders in the engine increases above 4, engine horsepower increases. there is negative correlation between mpg and horsepower.

In [25]:

```
plt.figure(figsize=(10,5))
ax = sns.countplot(x = 'origin', data = mpg_df)
ax.bar_label(ax.containers[0], label_type='edge')
plt.title("origin distribution", fontsize = 20)
plt.xlabel("origin", fontsize = 15)
plt.ylabel("car count", fontsize = 15)
plt.show()
```

We can see that USA alone produces more than 64% of the total cars in our dataset.

In [26]:

```
plt.figure(figsize=(10,5))
sns.pointplot(x = 'origin', y = 'mpg', data = mpg_df)
plt.title("origin against mpg", fontsize = 20)
plt.xlabel("origin", fontsize = 15)
plt.ylabel("mpg", fontsize = 15)
plt.show()
```

Although USA has the biggest count of our dataset, it produces cars we relatively very low mpg compared to Asia and Europe Asia is the leading contry in producing cars with high mpg with a mean close to 30.

In [27]:

```
plt.figure(figsize=(10,5))
sns.pointplot(x = 'origin', y = 'horsepower', data = mpg_df)
plt.title("origin against horsepower", fontsize = 20)
plt.xlabel("origin", fontsize = 15)
```

```
plt.ylabel("horsepower", fontsize = 15)
plt.show()
```

As expected, USA has the highest engine horsepower. This implies the previously observed conclusion that horsepower and mpg has a negative correlation.

In [28]:

```
plt.figure(figsize=(10,5))
sns.pointplot(x = 'model year', y = 'mpg', hue = 'origin', data = mpg_df, ci = None);
plt.title("model year against mpg with respect to origin", fontsize = 20)
plt.xlabel("model year", fontsize = 15)
plt.ylabel("mpg", fontsize = 15)
plt.show()
```

```
C:\Users\pandeyunny2315\AppData\Local\Temp\ipykernel_10128\2067465216.py:2: FutureWarning:
The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.
```

```
sns.pointplot(x = 'model year', y = 'mpg', hue = 'origin', data = mpg_df, ci = None);
```

This chart also implies the positive correlation between model year and mpg, and shows that Asia has been leading country in this industry, followed by Europe and USA comes in the last place.

In [29]:

```
plt.figure(figsize=(10,5))
sns.histplot(x = 'weight', data = mpg_df, color = '#4287f5')
plt.title("car weight distribution", fontsize = 20)
plt.xlabel("car weight", fontsize = 15)
plt.ylabel("car count", fontsize = 15)
plt.show()
```

car weight varies from 1500 to 5000, with the majority of cars at 2000.

In [30]:

```
plt.figure(figsize=(10,5))
ax = sns.barplot(x = 'model year', y = 'weight', data = mpg_df,
                 color = '#4287f5', ci = None)
plt.title("weight against model year", fontsize = 20)
plt.xlabel("model year", fontsize = 15)
plt.ylabel("car weight", fontsize = 15)
plt.show()
```

```
C:\Users\pandeyunny2315\AppData\Local\Temp\ipykernel_10128\1031872803.py:2: FutureWarning:
The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.
```

```
ax = sns.barplot(x = 'model year', y = 'weight', data = mpg_df,
```

In [31]:

```
plt.figure(figsize=(10,5))
plt.scatter(x = 'mpg', y = 'weight', data = mpg_df)
plt.title("weight against mpg", fontsize = 20)
plt.xlabel("mpg", fontsize = 15)
plt.ylabel("car weight", fontsize = 15)
plt.show()
```

mpg inceases as weight decreases over time, that indecates a stron correlation between them.

In [32]:

```
plt.figure(figsize=(10,5))
sns.barplot(x = 'origin', y = 'weight', data = mpg_df)
plt.title("weight against origin", fontsize = 20)
```

```
plt.xlabel("origin", fontsize = 15)
plt.ylabel("car weight", fontsize = 15)
plt.show()
```

As expected, USA gets the highest possible weight values, that explains a lot about its low mpg and high horsepower. Asia produces the lightest cars compared to USA and Europe, this also explains its high mpg and low horsepower.

In [33]:

```
plt.figure(figsize=(10,5))
plt.scatter(x = 'horsepower', y = 'acceleration', data = mpg_df)
plt.title("horsepower against acceleration", fontsize = 20)
plt.xlabel("horsepower", fontsize = 15)
plt.ylabel("acceleration", fontsize = 15)
plt.show()
```

Conclusion¶

In this section, we'd add the conclusions we draw from the previous visualisations.

As years pass after 1973, there has been a noticeable increase in mpg.

As cylinders in the engine increases above 4, MPG decreases and engine horsepower increases.

That indicates negative correlation between mpg and horsepower.

mpg increases as weight decreases over time, that also indicates a strong correlation between them.

Although USA has the biggest count of produced cars, its cars have relatively very low mpg, thus the highest possible weight compared to Asia and Europe.

Asia is the leading country in producing cars with high mpg with a mean close to 30, and it produces the lightest cars. We can spot a negative correlation between acceleration and horsepower, this means that it has a positive one with mpg.

In []: